

# ML Online Course

## Linear Regression with multiple Variables

# Linear Regression with multiple variables

$x_j^{(i)}$  = value of feature  $j$  in  $i^{\text{th}}$  training example.

$$\begin{bmatrix} x_1^1 \\ x_1^2 \\ x_1^3 \end{bmatrix}$$

$x^{(i)}$  = input (features) of  $i^{\text{th}}$  training example.

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_3 x_3 + \dots + \theta_n x_n$$

For convenience of notation, define  $x_0 = 1$  ( $x_0^{(i)} = 1$ )

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n \quad [\theta_0 \theta_1 \dots \theta_n] \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

$$= \theta^T x$$

Multivariate Linear Regression

$$X = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} \\ x_0^{(2)} & x_1^{(2)} \\ x_0^{(3)} & x_1^{(3)} \end{bmatrix}, \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} \quad h(X) = X\theta$$

and  $x_0^{(i)} = 1$

## Gradient Descent

Hypothesis:  $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$

Parameters:  $\theta$   $n+1$ -dimensional vector

Cost function:  $J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Gradient descent:  $J(\theta)$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}  
 $\uparrow$  (simultaneously update for every  $j = 0, \dots, n$ )

For  $\theta$  ( $n \geq 1$ ):

Repeat {

$$\theta_j := \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \right]$$

} (simultaneously update  $\theta_j$  for  $j=0, \dots, n$ )

Special case:  $\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_0^{(i)}$

Feature Skills: Feature Scaling

Tricks to converge well

Idea: Make sure features are on a similar scale

(in order to make GD converge more quickly)

如果 scale 不 similar

那么 Contour 会特别的扁

这样的话 converge 特别的慢

on a similar scale it's

contour 接近一个圆



Feature scaling: Dividing the input values by the range of input variable, resulting in a new range of just 1. Get every feature into approximately a  $-1 \leq x_i \leq 1$  range.

Mean normalization.

Replace  $x_i$  with  $(x_i - \mu_i)$  to make features have approximately zero mean.

(Do not apply to  $x_0 = 1$ )

$x_1 \leftarrow x_1 - \mu_1$  — avg value of  $x_1$  in training set

$$x_1 \leftarrow \frac{x_1 - \mu_1}{s_1}$$

range (max-min)  
standard deviation

$$x_2 \leftarrow \frac{x_2 - \mu_2}{s_2}$$

Matrix Notation

$$\theta := \theta - \alpha \nabla J(\theta)$$

$$\nabla J(\theta) = \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_0} \\ \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{bmatrix}$$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_j} &= \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)} \\ &= \frac{1}{m} \sum_{i=1}^m x_j^{(i)} \cdot (h_\theta(x^{(i)}) - y^{(i)}) \end{aligned}$$

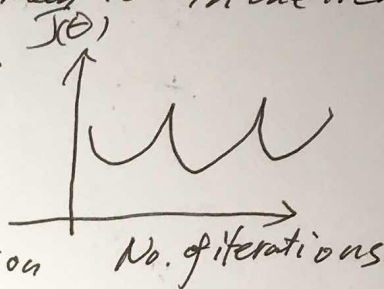
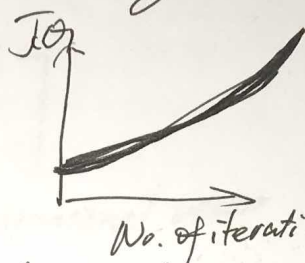
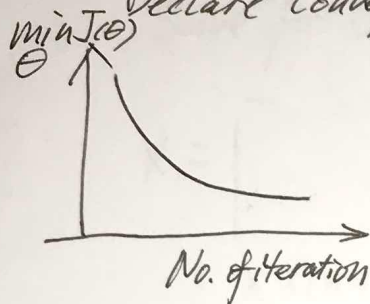
$$= \frac{1}{m} \vec{x}_j^T (X\theta - \vec{y})$$

$$\nabla J(\theta) = \frac{1}{m} X^T (X\theta - \vec{y}) \Rightarrow \theta := \theta - \frac{\alpha}{m} X^T (X\theta - \vec{y})$$

- "Debugging": How to make sure gradient descent is working correctly
- How to choose learning rate  $\alpha$

Convergence test:

Declare Convergence if  $J(\theta)$  decrease by less than  $10^{-3}$  in one iteration.



For sufficient small  $\alpha$ ,  $J(\theta)$  should decrease on every iteration

But if  $\alpha$  is too small, GD can be slow to converge

Summary:

- If  $\alpha$  is too small: slow convergence

If  $\alpha$  is too large:  $J(\theta)$  may not decrease on every iteration; may not converge.

To choose  $\alpha$ , try

..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1

Polynomial Regression:

Our hypothesis function need not be linear if it does not fit the data well.

Feature scaling would be important in the case when we use GD.



Normal equation: Method to solve for  $\theta$  analytically

No Need to do Feature Scaling

$$\theta \in \mathbb{R}^{n+1} \quad J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \dots \stackrel{\text{set}}{=} 0 \text{ (for every } j \text{)}$$

Solve for  $\theta_0, \theta_1, \dots, \theta_n$

$$X = \begin{bmatrix} \quad \end{bmatrix} \quad y = \begin{bmatrix} \quad \end{bmatrix}$$

$X$  is  $(n+1) \times m$  dimensional vector

$$\theta = (X^T X)^{-1} X^T y$$

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \text{---} (x^{(2)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix}$$

(design matrix)  $m \times (n+1)$

For example.

$$x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & x_1^{(1)} \\ \vdots & \vdots \\ 1 & x_1^{(m)} \end{bmatrix} \quad y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

$$\theta = (X^T X)^{-1} X^T y \quad \text{Octave: pinv}(X' * X) * X' * y$$

$m$  training examples,  $n$  features

GD  
Need to choose  $\alpha$   
Needs many iterations  
Works well even when  $n$  is large  
 $O(kn^2)$

Normal Equation

~~Do not~~ need to choose  $\alpha$

Don't need to iterate  
Need to compute  $(X^T X)^{-1}$   $O(n^3)$   
Slow if  $n$  is very large

Normal equation and non-invertibility

$$\theta = (X^T X)^{-1} X^T y$$

what if  $X^T X$  is non-invertible?

- Redundant features
- Too many features (e.g.  $m \leq n$ )
- Delete some features, or use regularization.

↑  
↓  
← →  
\* .^ ./

Matlab

$V = \text{linspace}(0, 3, 8)$

$V = 2:0.2:3$

$V = C'$

~~plot~~  $\text{plot}(x, y)$

'm-s' magenta, dotted line with square markers  
'g--\*' green dashed line with star markers

'r-' red, solid line with no markers

$x\text{label}('time [s]')$

$y\text{label}('amplitude')$

$\text{title}('my plot')$

$\text{legend}('y(t)')$

grid  
annotation

matrices

$A = [-4, 1.9, -3.2, -12; -0.25, 2, 9, 0.3; 0.1, 7, -1, 5]$

$\begin{bmatrix} -4 & 1.9 & -3.2 & -12 \\ -0.25 & 2 & 9 & 0.3 \\ 0.1 & 7 & -1 & 5 \end{bmatrix}$

Array Creation

rand

ones

eye

randi

randn

zeros

toeplitz

vander

diag

magic

$\text{randi}([0, 2], 3, 4)$  random integers between 0 and 2

normal random variables

$\text{zeros}(10, 3)$

vandermonde matrix

hilb

$$S_1 = M(3, 2)$$

$$S_2 = M([2, 3], 2)$$

$$S_3 = M(2:4, 2)$$

$$S_4 = M(:, 2)$$

u vector  
 $L_1 = \text{length}(u)$

length(m) will return the max of columns

matrix  
 $s = \text{size}(m)$

$C = [A, B]$  horizontal concatenating

$C = [A; B]$  vertical concatenating

$A * x$

matrix multiplication

$\sim$  = not equal to

vector  
 $I = V < 0.05$

$V(I) = 0$

for  $n = 1:6$

$y(n+1) = y(n) - 0.1 * y(n)$

end

IF

Else

End

while

end



1=2 % false

1~2

1&2 0 % And

1||0 % or

Xor(1,0)

who

whos

save hello.mat

~~hello~~ save hello.mat v  
save hello.tex v -ascii

A(3,2)

A(2,:)

A(:,2)

A([1,3],:)

a=3; semicolon suppressing count  
format(long)

format(short)

V=[1 2 3]

V=[1;2;3]

V=1:0.1:2

1→2 0.1-间隔

1:6

1→6 1间隔

ones(2,3)

zeros(1,3)

rand(1,3)

rand(3,3)

randn(1,3) Gaussian Distribution

hist(w)

eye()

~~help eye~~ help eye

size(A) size(A,1) size(A,2)

V=[1 2 3 4]

length(V) ans=4

length(A) ans=3

length([1,2,3,4,5])

pwd cd ls

load filename.dat

load('\_.dat')